


# Operating Systems 1: Processes

## Introduction

- We are using an OS daily. Even currently while watching the class.
- [POLL] Which OS are you using right now?
- Why learn OS?
  - To understand how the applications we develop or use really work. How are they able to get things done?

## What is an OS?

### Resource Manager



- Tell how every process basically needs different resources to do their work and to avoid conflicts, there needs to be someone to manage them. Like a Program Manager in a company.

### Functionalities for Developers

- Interacting with Network
- Working with File System
- Imagine having to write the code to rotate the hard disk to fetch a file at a particular location.

## Uni v/s Multiprogramming Computers

- Multiple programs being executed at a time (not necessary concurrently) instead of a single.
- Run the other one when first is waiting for some resources.
- Make full use of CPU
- [Show multiple programs running in parallel on Mac]
- Multitasking is the same as multiprogramming. Just used in windows.

## Types of Multiprogramming

- Single v/s Multi User
- Preemptive v/s Non Preemptive -> OS can force process to be removed even before completion

## fork()

- System Call
- Makes CPU come in Kernel Mode
- What is Kernel Mode
  - Non Preemptive
  - Key OS functionality
  - Atomic Execution
- Creates a new process with copy of the current variables

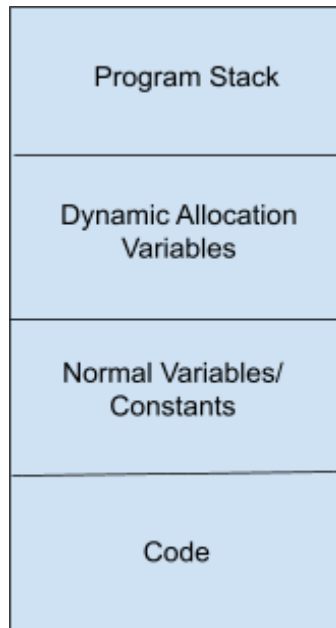
```
int main() {  
    int a = 2;  
    int b = 3;  
    id = fork();  
    print("hi");  
}
```

- Id = zero in the child process.
- Non zero in the parent process.
- Less than zero means some error.
- [Show demo of how the program will work with 1 and 2 fork]
- [Show that n forks mean  $2^n$  processes]
- How OS works
  - Even OS is a process that is started at the start of a program by BIOS that internally would start other processes. [Show systemd in Mac]

## Process

- Program in execution
- Has resources allocated to it
- Unit of execution

What a process consists of



## Process in memory

[Give hints regarding what all details wrt a process might be needed to store by an OS.  
Assume process class is there.]

- ID
- CPU related info (Priority etc.)
- Memory info (Limits etc.)
- Opened Files
- I/O related details
- Protection (Kernel/ User Mode)

All of these go in PCB (Process Control Block).

## CPU Scheduling

### CPU Bound v/s I/O Bound Processes

- I/O bound like Interactive programs. Most of the time waiting for I/O than using CPU.
- CPU Bound like scientific applications doing calculations

## How I/O Happens

### Interrupts

- Program raises a CPU Interrupt telling CPU that it is waiting for some I/O to happen
- Interrupts are signals emitted to tell the processor that some interruption is required.
- CPU sits idle till the I/O (network call, command line input etc.) which is very slow wrt CPU completes.

- This is a waste of CPU time and CPU can instead spend this time doing something useful.

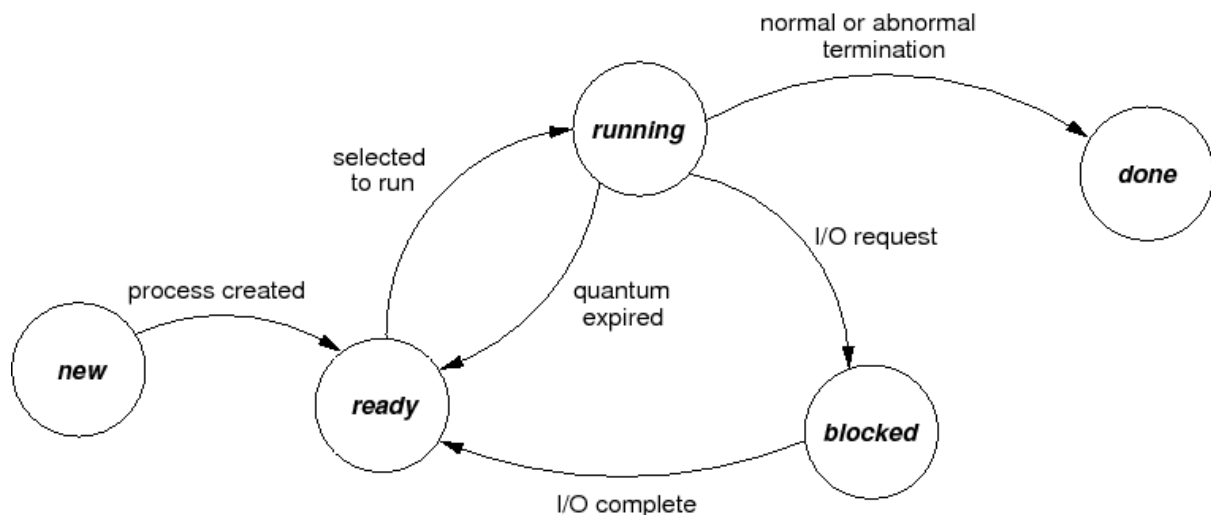
## Why CPU Scheduling

- Utilize the CPU to fullest

### Metrics wrt Processes

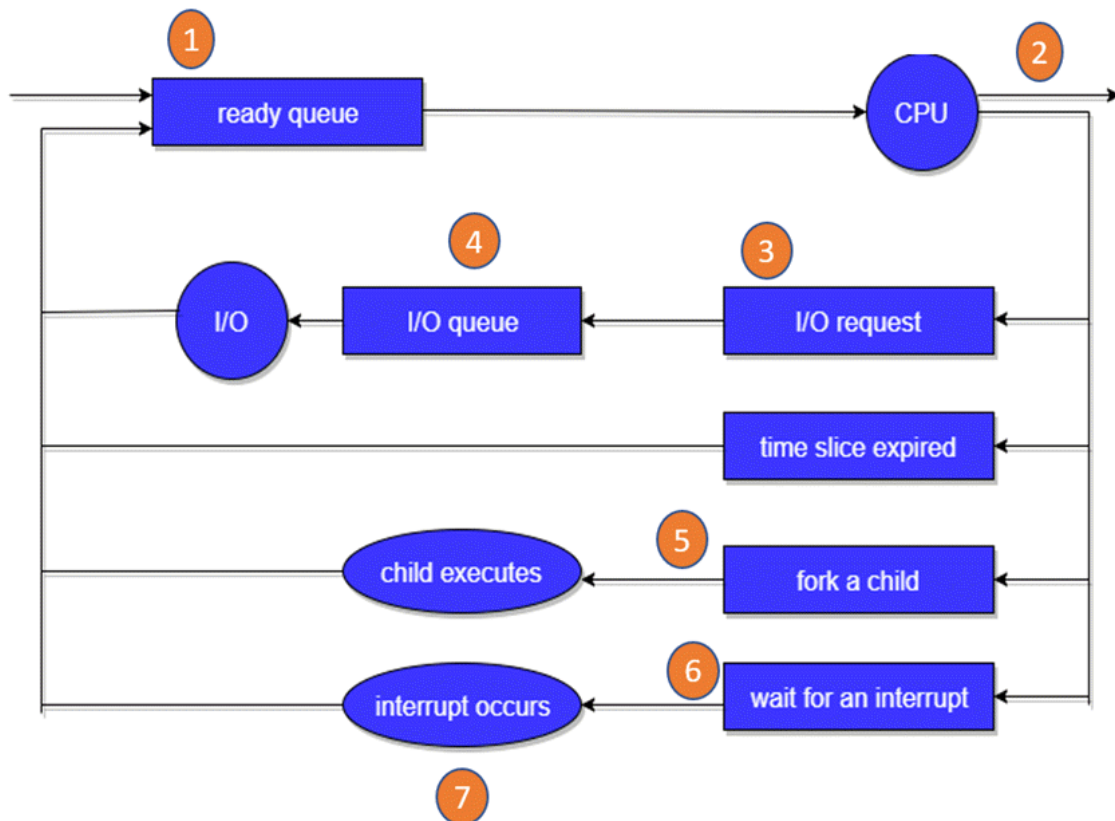
- Throughput: Number of process per unit time
- Arrival Time: When a process comes to CPU
- Wait Time: Time till completion when process isn't executing but is waiting for I/O or waiting to be assigned to CPU.
- Burst Time: Time needed to make the process execute completely.
- Completion Time: Time when a process completes its execution,
- I/O Burst Wait Time: Time for I/O to happen
- Turnaround Time = Completion Time - Arrival Time
- Response Time = Time when process first starts running (gets assigned to CPU)
- Schedule Time = Time to completely finish all processes
- Deadline: Expectation by which process should ideally finish
  - Underrun: When process finishes before deadline
  - Overrun: When process finishes after deadline

## Process Lifecycle



## Process Scheduling Queues

When a process is not running, it is waiting in one of the scheduling queues.



### Dispatcher

- Picks the process from Ready Queue and assigns it to the CPU for scheduling.
- Moves the PCB out of CPU into the registers for execution.

### Context Switch

- We as humans when we do multiple tasks in parallel, it takes some time to recall what had happened earlier.
- Moving a process out of register, storing it in memory in ready queue, replacing it with another process is known as context switch
- Context switch takes time and thus an algorithm that involves a lot of context switch isn't good.

### Scheduling Algorithms

- Assume no I/O
- No context switch time

### FCFS

- Non Preemptive.
- When we have to select a new process from the ready queue, select the process that has arrived the earliest.

| PId | Arrival Time | Burst Time | Response Time | Completion Time | TAT | Wait Time |
|-----|--------------|------------|---------------|-----------------|-----|-----------|
| 1   | 0            | 4          |               |                 |     |           |
| 2   | 1            | 5          |               |                 |     |           |
| 3   | 2            | 6          |               |                 |     |           |
| 4   | 3            | 8          |               |                 |     |           |
| 5   | 4            | 2          |               |                 |     |           |
| 6   | 5            | 4          |               |                 |     |           |

[Show how scheduling will work for above by filling the above table, as well as via a Gantt chart]

### SRTF

- Shortest Response Time First
- Preemptive

| PId | Arrival Time | Burst Time |
|-----|--------------|------------|
| 1   | 0            | 8          |
| 2   | 1            | 6          |
| 3   | 2            | 4          |
| 4   | 3            | 2          |
| 5   | 4            | 6          |
| 6   | 5            | 1          |

[Show the scheduling via Gantt chart and show how preemptive behaviour happens.]

### Starvation

- Some processes have to keep on waiting for very long and don't get CPU time because of other new processes coming ahead.

### Round Robin

- Preemptive
- Time Quantum (q)
- After every q units, check the process most ahead in the queue and schedule that.

- New process gets added behind the queue.
- Current process gets added at the end when preempted.

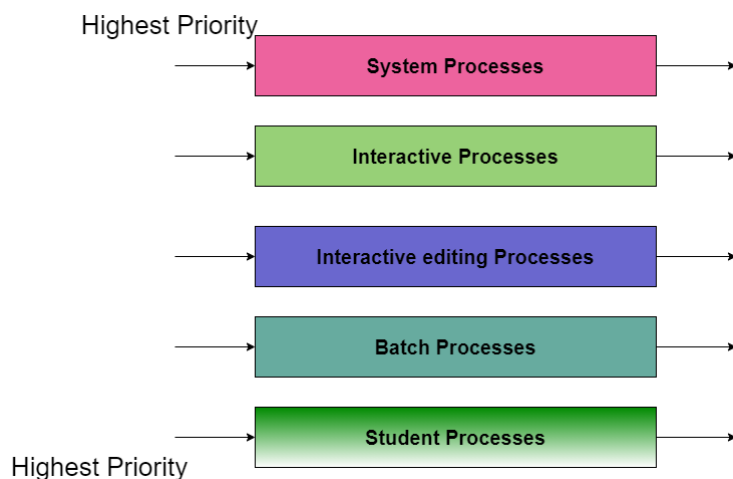
| PId | Arrival Time | Burst Time |
|-----|--------------|------------|
| 1   | 0            | 4          |
| 2   | 1            | 5          |
| 3   | 2            | 6          |
| 4   | 3            | 3          |
| 5   | 4            | 4          |
| 6   | 5            | 1          |

### Homework

- Learn about SJF and Priority scheduling.
- [Give a brief idea about them. The only difference is in one, the process with highest priority is selected, in the other the process with the shortest time.]

### Multilevel Queues

- In real systems, a single scheduling algorithm might not work well.
- For different type of processes, we might need different scheduling algorithms.
- Example:
  - For I/O bound processes, RR with short time
  - For CPU intensive, maybe something else.



- If there is a process in the highest queue, pick that. Else go to lower queues. Each queue internally may have its own scheduling algorithm.